

# E-Commerce Sales Hive Case Study

## Contents

Problem Statement:.....	2
Objective .....	2
Steps Involved .....	2
Create EMR and Load the data to HDFS for Analysis.....	2
Create Key Pair .....	2
Create EMR Cluster .....	3
Create S3 Bucket and upload the Data .....	4
Enable SSH Connection and Connect Putty .....	5
Setup the directory .....	5
Starting Hive:.....	7
Applying Optimization Techniques Columns Partitioning and Bucketing: .....	9
Questions and Queries to Answers:.....	14
Dropping the Database and Terminating the Cluster .....	27
Drop Database: .....	27
Terminate EMR: .....	27

## Problem Statement:

A tech companies is exploring ways to improve their sales. They want to start by analysing customer behaviour and gaining insights about product trends.

Here, the role of big data analysts is among the most sought-after to gain the insights from abundance of data.

## Objective

Through this assignment, as a big data analyst, we will extract data and gather insights from a real-life data set of an e-commerce company. We will analyse and gain insights about the clickstream data from a website so that we can extract insights about the customers behaviour.

## Steps Involved

The steps involved in the entire process are as follows:

- ✓ Create the Key pair in AWS
- ✓ Create S3 bucket and copy the csv data files into s3 bucket
- ✓ Create the EMR cluster and launch EMR that utilizes the Hive services
- ✓ Move the data from the S3 bucket into the HDFS
- ✓ Creating the database and launching Hive queries on EMR cluster
- ✓ Creating the structure of database
- ✓ Use partitioning and bucketing to run your queries as efficiently as possible
- ✓ Run Hive queries to verify the analysis on given dataset.
- ✓ Drop the database, and
- ✓ Terminate the cluster

## Create EMR and Load the data to HDFS for Analysis

### Create Key Pair

- Open EC2 Management Console and select Create Key Pair. Enter a name for Key Pair, keeping the type as “RSA” amd file format as “.ppk” and Select Create Key Pair. The Key Pair is downloaded.

Create key pair [Info](#)


**Key pair**  
A key pair, consisting of a private key and a public key, is a set of security credentials that you use to prove your identity when connecting to an instance.

Name  
  
The name can include up to 255 ASCII characters. It can't include leading or trailing spaces.

Key pair type [Info](#)  
☒ RSA  
☐ ED25519

Private key file format  
☐ .pem  
For use with OpenSSH  
☒ .ppk  
For use with PuTTY

- A new key pair “Hive\_Assignment\_K” has been created to be used for this case study.



## Create EMR Cluster

- Go to EMR > Select Create Cluster. Selected the General Configuration, Software Configuration, and Hardware Configuration as shown in below images. Select “Hive\_Assignment\_K” in EC2 key pair. Click “Create Cluster”.

Amazon EMR

EMR Studio

Create cluster View details Clone Terminate

Filter: All clusters Filter clusters ... 15 clusters (all loaded)

### Create Cluster - Quick Options [Go to advanced options](#)

#### General Configuration

Cluster name

☐ Logging

Launch mode ☒ Cluster ☐ Step execution

#### Software configuration

Release

Applications

- ☒ Core Hadoop: Hadoop 2.8.5 with Ganglia 3.7.2, Hive 2.3.6, Hue 4.4.0, Mahout 0.13.0, Pig 0.17.0, and Tez 0.9.2
- ☐ HBase: HBase 1.4.10 with Ganglia 3.7.2, Hadoop 2.8.5, Hive 2.3.6, Hue 4.4.0, Phoenix 4.14.3, and ZooKeeper 3.4.14
- ☐ Presto: Presto 0.227 with Hadoop 2.8.5 HDFS and Hive 2.3.6 Metastore
- ☐ Spark: Spark 2.4.4 on Hadoop 2.8.5 YARN with Ganglia 3.7.2 and Zeppelin 0.8.2
- ☐ Use AWS Glue Data Catalog for table metadata

#### Hardware configuration

Instance type  The selected instance type adds 32 GiB of GP2 EBS storage per instance by default. [Learn more](#)

Number of instances  (1 master and 1 core nodes)

#### Security and access

EC2 key pair  [Learn how to create an EC2 key pair.](#)

Permissions ☒ Default ☐ Custom

Use default IAM roles. If roles are not present, they will be automatically created for you with managed policies for automatic policy updates.

EMR role [EMR\\_DefaultRole](#) ☐ Use EMR\_DefaultRole\_V2

EC2 instance profile [EMR\\_EC2\\_DefaultRole](#)

Cancel Create cluster

- The EMR Cluster creation and configuration begins in the backend. The **Starting** symbol beside Cluster name indicates it.

The screenshot shows the AWS Management Console interface for an Amazon EMR cluster. The cluster is named 'Hive\_Case\_Study' and is in the 'Starting' state. The console displays various tabs for the cluster, including Summary, Application user interfaces, Monitoring, Hardware, Configurations, Events, Steps, and Bootstrap actions. The Summary tab is selected, showing the cluster ID (j-381DZM4AD40PU), creation date (2022-02-01 12:00 (UTC+5:30)), and elapsed time (0 seconds). It also shows the termination protection status (Off) and the master public DNS. The Configuration details section lists the release label (emr-5.29.0), Hadoop distribution (Amazon 2.8.5), and the applications installed (Ganglia 3.7.2, Hive 2.3.6, Hue 4.4.0, Mahout 0.13.0, Pig 0.17.0, Tez 0.9.2). A warning message at the top indicates that auto-termination is not available for this account when using this release of EMR.

Filter: Active clusters Filter clusters ... 1 cluster (all loaded)

Name	ID	Status	Creation time (UTC+5:30)	Elapsed time
<a href="#">Hive_Case_Study</a>	j-381DZM4AD40PU	Starting	2022-02-01 12:00 (UTC+5:30)	1 minute

## Create S3 Bucket and upload the Data

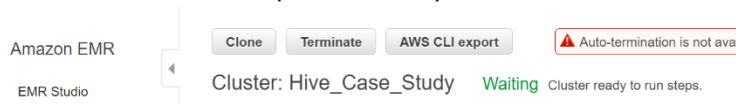
- Go to Amazon S3 > Select Create Bucket > Enter Bucket Name “[hive.casestudy](#)” > Unselect “Block all public access” > Select Create Bucket.

The screenshot shows the Amazon S3 console interface for the 'hive.casestudy' bucket. The console displays the bucket name, a search bar, and a list of objects. The 'Objects' tab is selected, showing a list of objects with columns for Name, Type, Last modified, Size, and Storage class. The 'Upload' button is visible, indicating that files can be uploaded to the bucket.


- Upload the data files “2019-Oct” and “2019-Nov” that are provided by UpGrad.
- After the files are uploaded, Copy S3 URL of both the uploaded files:
  - s3://hive.casestudy/2019-Oct.csv
  - s3://hive.casestudy/2019-Nov.csv

### Enable SSH Connection and Connect Putty

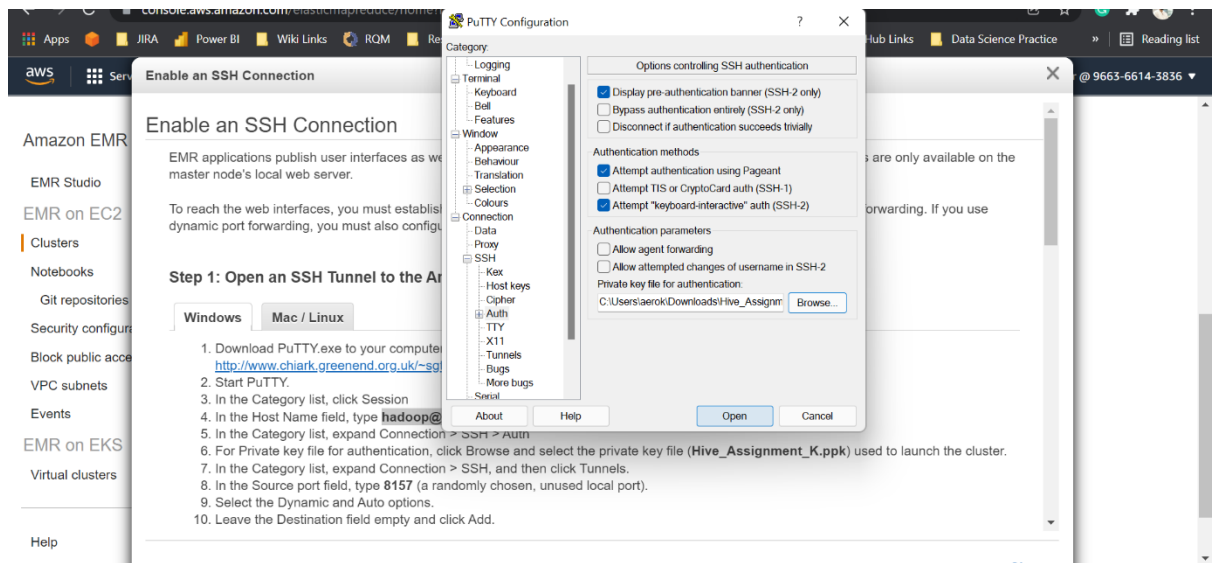
- Go to Amazon EMR > Open the newly created active cluster after it is in **waiting** >



- Scroll down to Application User > Select **Enable an SSH Connection** > Copy the Host Name Field **"hadoop@ec2-54-196-3-223.compute-1.amazonaws.com"**



- Open Putty > Enter the host name [hadoop@ec2-54-196-3-223.compute-1.amazonaws.com](https://hadoop@ec2-54-196-3-223.compute-1.amazonaws.com) and select the key pair.
- Once cluster in running state we have to click on Master public DNS. Open the putty configuration and then give the host name (master node DNS) and then browse to the private key file location by clicking on Connection → SSH → Auth. Now we need to open Putty and connect to the master node by selecting the .ppk file.



### Setup the directory

- EMR Opens, and below screen reflects that Connection to Hadoop is successful:

[illegible]

- Check the existing directories, and create a directory named 'Upgrad\_Hive\_Case\_Study' in Hadoop. > 'Upgrad\_Hive\_Case\_Study' is now created in Hadoop:

**Command:**

```
hadoop fs -ls /
```

```
hadoop fs -mkdir /Upgrad_Hive_Case_Study
```

```
hadoop fs -ls /
```

```
[hadoop@ip-172-31-44-242 ~]$ hadoop fs -ls /
Found 4 items
drwxr-xr-x   - hdfs  hadoop          0 2022-02-02 06:24 /apps
drwxrwxrwt   - hdfs  hadoop          0 2022-02-02 06:26 /tmp
drwxr-xr-x   - hdfs  hadoop          0 2022-02-02 06:24 /user
drwxr-xr-x   - hdfs  hadoop          0 2022-02-02 06:24 /var
[hadoop@ip-172-31-44-242 ~]$ hadoop fs -mkdir /Upgrad_Hive_Case_Study
[hadoop@ip-172-31-44-242 ~]$ hadoop fs -ls /
Found 5 items
drwxr-xr-x   - hadoop hadoop          0 2022-02-02 06:29 /Upgrad_Hive_Case_Study
drwxr-xr-x   - hdfs  hadoop          0 2022-02-02 06:24 /apps
drwxrwxrwt   - hdfs  hadoop          0 2022-02-02 06:26 /tmp
drwxr-xr-x   - hdfs  hadoop          0 2022-02-02 06:24 /user
drwxr-xr-x   - hdfs  hadoop          0 2022-02-02 06:24 /var
[hadoop@ip-172-31-44-242 ~]$
```

- Move the data from the s3 buckets to the HDFS using the distributed copy command. Loading the s3 public data set to created directory “Upgrad\_Hive\_Case\_Study” in hadoop .

**Command:**

```
hadoop distcp 's3://hive.casestudy/2019-Oct.csv' /Upgrad_Hive_Case_Study/2019-Oct.csv
```

```
hadoop distcp 's3://hive.casestudy/2019-Nov.csv' /Upgrad Hive Case Study/2019-Nov.csv
```

```
[hadoop@ip-172-31-44-242 ~]$ hadoop distcp 's3://hive.casestudy/2019-Oct.csv' /Upgrad_Hive_Case_Study /2019-Oct.csv
22/02/02 06:32:34 INFO tools.DistCp: Input Options: DistCpOptions{atomicCommit=false, syncFolder=false, deleteMissing=false, ignoreFailures=false, overwrite=false, skipCRC=false, blockin
```

```

Map-Reduce Framework
  Map input records=2
  Map output records=0
  Input split bytes=272
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=630
  CPU time spent (ms)=21260
  Physical memory (bytes) snapshot=772878336
  Virtual memory (bytes) snapshot=6558760960
  Total committed heap usage (bytes)=632291328
File Input Format Counters
  Bytes Read=650
File Output Format Counters
  Bytes Written=0
DistCp Counters
  Bytes Copied=482542278
  Bytes Expected=482542278
  Files Copied=2

```

```

[hadoop@ip-172-31-44-242 ~]$ hadoop distcp 's3://hive.casestudy/2019-Nov.csv' /Upgrad_Hive_Ca
se_Study /2019-Nov.csv
22/02/02 06:35:55 INFO tools.DistCp: Input Options: DistCpOptions{atomicCommit=false, syncFol
der=false, deleteMissing=false, ignoreFailures=false, overwrite=false, skipCRC=false, blockin

```

```

Map-Reduce Framework
  Map input records=2
  Map output records=0
  Input split bytes=272
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=593
  CPU time spent (ms)=22090
  Physical memory (bytes) snapshot=804085760
  Virtual memory (bytes) snapshot=6560874496
  Total committed heap usage (bytes)=671612928
File Input Format Counters
  Bytes Read=650
File Output Format Counters
  Bytes Written=0
DistCp Counters
  Bytes Copied=545839412
  Bytes Expected=545839412
  Files Copied=2

```

### Starting Hive:

- It's time to Set the data in Hive,
- Launch Hive > show databases > CREATE a DATABASE "hive\_data" > Check if Database is created.

#### Codes:

Hive

```
CREATE DATABASE IF NOT EXISTS hive_data;
```

```
SHOW DATABASES;
```

```
DESCRIBE DATABASE hive_data;
```

```
[hadoop@ip-172-31-44-242 ~]$ hive
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false
hive> CREATE DATABASE IF NOT EXISTS hive_data;
OK
Time taken: 1.019 seconds
hive> SHOW DATABASES;
OK
default
hive_data
Time taken: 0.183 seconds, Fetched: 2 row(s)
hive> DESCRIBE DATABASE hive_data;
OK
hive_data          hdfs://ip-172-31-44-242.ec2.internal:8020/user/hive/warehouse/hive_data.db
hive>
```

Use the database and create an external table “RetailDB”:

**Query:**

use hive\_data;

CREATE EXTERNAL TABLE IF NOT EXISTS RetailDB (event\_time timestamp, event\_type string, product\_id string, category\_id string, category\_code string, brand string, price decimal(10,3), user\_id bigint, user\_session string) ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' WITH SERDEPROPERTIES ("separatorChar" = ",", "quoteChar" = "\"", "escapeChar" = "\\") stored as textfile

DESCRIBE RetailDB;

```
hive> use hive_data;
OK
Time taken: 0.026 seconds
hive> CREATE EXTERNAL TABLE IF NOT EXISTS RetailDB (event_time timestamp, event_type string, product_id string, category_id string, category_code string, brand string, price decimal(10,3), user_id bigint, user_session string) ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' WITH SERDEPROPERTIES ("separatorChar" = ",", "quoteChar" = "\"", "escapeChar" = "\\") stored as textfile;
OK
Time taken: 1.079 seconds
hive> DESCRIBE RetailDB;
OK
event_time          string              from deserializer
event_type          string              from deserializer
product_id          string              from deserializer
category_id         string              from deserializer
category_code       string              from deserializer
brand               string              from deserializer
price               string              from deserializer
user_id             string              from deserializer
user_session        string              from deserializer
Time taken: 0.137 seconds, Fetched: 9 row(s)
```

- Load the input Data into “RetailDB” table:

**Query:**

LOAD DATA INPATH 's3://hive.casestudy/2019-Oct.csv' INTO TABLE RetailDB;

LOAD DATA INPATH 's3://hive.casestudy/2019-Nov.csv' INTO TABLE RetailDB;



```
hive> LOAD DATA INPATH 's3://hive.casestudy/2019-Oct.csv' INTO TABLE RetailDB;
Loading data to table hive_data.retaildb
OK
Time taken: 17.997 seconds
hive> LOAD DATA INPATH 's3://hive.casestudy/2019-Nov.csv' INTO TABLE RetailDB;
Loading data to table hive_data.retaildb
OK
Time taken: 10.8 seconds
hive>
```

- Test the loaded data and check if data of November and October are loaded:

#### Query:

```
SELECT * FROM RetailDB WHERE MONTH(event_time)=11 limit 5;
```

```
SELECT * FROM RetailDB WHERE MONTH(event_time)=10 limit 5;
```

```
hive> SELECT * FROM RetailDB WHERE MONTH(event_time)=11 limit 5;
OK
2019-11-01 00:00:02 UTC view      5802432 1487580009286598681      0.32      56207
6640      09fafd6c-6c99-46b1-834f-33527f4de241
2019-11-01 00:00:09 UTC cart      5844397 1487580006317032337      2.38      55332
9724      2067216c-31b5-455d-alcc-af0575a34ffb
2019-11-01 00:00:10 UTC view      5837166 1783999064103190764      pnb      22.22      55613
8645      57ed222e-a54a-4907-9944-5a875c2d7f4f
2019-11-01 00:00:11 UTC cart      5876812 1487580010100293687      jessnail      3.165
64506666      186c1951-8052-4b37-adce-dd9644b1d5f7
2019-11-01 00:00:24 UTC remove_from_cart      5826182 1487580007483048900      3
.33      553329724      2067216c-31b5-455d-alcc-af0575a34ffb
Time taken: 0.802 seconds, Fetched: 5 row(s)
hive> ;
hive> ;
hive> SELECT * FROM RetailDB WHERE MONTH(event_time)=10 limit 5;
OK
2019-10-01 00:00:00 UTC cart      5773203 1487580005134238553      runail      2.62      46324
0011      26dd6e6e-4dac-4778-8d2c-92e149dab885
2019-10-01 00:00:03 UTC cart      5773353 1487580005134238553      runail      2.62      46324
0011      26dd6e6e-4dac-4778-8d2c-92e149dab885
2019-10-01 00:00:07 UTC cart      5881589 2151191071051219817      lovely      13.48      42968
1830      49e8d843-adf3-428b-a2c3-fe8bc6a307c9
2019-10-01 00:00:07 UTC cart      5723490 1487580005134238553      runail      2.62      46324
0011      26dd6e6e-4dac-4778-8d2c-92e149dab885
2019-10-01 00:00:15 UTC cart      5881449 1487580013522845895      lovely      0.56      42968
1830      49e8d843-adf3-428b-a2c3-fe8bc6a307c9
Time taken: 0.29 seconds, Fetched: 5 row(s)
hive>
```

Further, randomly check if query works for aggregate functions:

```
hive> SELECT Avg(price) AS Oct_Revenue FROM RetailDB WHERE MONTH(event_time) = '10' AND event
_type = 'purchase';
Query ID = hadoop_20220202070346_b6bb6393-b403-4163-8149-f69e7950b66e
Total jobs = 1
Launching Job 1 out of 1
```

```
Map 1: 8/8      Reducer 2: 1/1
OK
4.932492061036307
Time taken: 92.006 seconds, Fetched: 1 row(s)
```

#### Applying Optimization Techniques Columns Partitioning and Bucketing:

- Start with enabling the dynamic partitioning and bucketing:

**Code:**

```
set hive.exec.dynamic.partition.mode = nonstrict;

set hive.exec.dynamic.partition = true;

set hive.enforce.bucketing = true;
```

```
hive> set hive.exec.dynamic.partition.mode = nonstrict;
hive> set hive.exec.dynamic.partition = true;
hive> set hive.enforce.bucketing = true;
```

Let's create and test two different combinations of partitioning and bucketing:

- Create a Table "RetailDB\_EU1", PARTITIONED BY(event\_type string) CLUSTERED BY (user\_id) INTO 10 buckets
- Create a Table "RetailDB\_EC2", PARTITIONED BY(event\_type string) CLUSTERED BY (category\_id) INTO 10 buckets

**Query:**

```
CREATE TABLE IF NOT EXISTS RetailDB_EU1 (event_time timestamp, product_id string, category_id string,
category_code string, brand string, price float, user_id bigint, user_session string) PARTITIONED BY
(event_type string) CLUSTERED BY (user_id) INTO 10 BUCKETS ROW FORMAT SERDE
'org.apache.hadoop.hive.serde2.OpenCSVSerde' STORED AS TEXTFILE;
```

```
CREATE TABLE IF NOT EXISTS RetailDB_EC2 (event_time timestamp, product_id string, category_id string,
category_code string, brand string, price float, user_id bigint, user_session string) PARTITIONED BY
(event_type string) CLUSTERED BY (category_id) INTO 10 BUCKETS ROW FORMAT SERDE
'org.apache.hadoop.hive.serde2.OpenCSVSerde' STORED AS TEXTFILE;
```

```
hive> CREATE TABLE IF NOT EXISTS RetailDB_EU1 (event_time timestamp, product_id string, category_id string,
category_code string, brand string, price float, user_id bigint, user_session string) PARTITIONED BY (event_type string) CLUSTERED BY (user_id) INTO 10 BUCKETS ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' STORED AS TEXTFILE;
OK
Time taken: 0.101 seconds
hive> CREATE TABLE IF NOT EXISTS RetailDB_EC2 (event_time timestamp, product_id string, category_id string,
category_code string, brand string, price float, user_id bigint, user_session string) PARTITIONED BY (event_type string) CLUSTERED BY (category_id) INTO 10 BUCKETS ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' STORED AS TEXTFILE;
OK
Time taken: 0.093 seconds
```

Check Properties of both Tables:

**Query:**

```
Describe RetailDB_EU1;
```

```
Describe RetailDB_EC2;
```

```

hive> Describe RetailDB_EU1;
OK
event_time          string          from deserializer
product_id          string          from deserializer
category_id         string          from deserializer
category_code       string          from deserializer
brand               string          from deserializer
price               string          from deserializer
user_id             string          from deserializer
user_session        string          from deserializer
event_type          string

# Partition Information
# col_name          data_type      comment

event_type          string
Time taken: 0.159 seconds, Fetched: 14 row(s)
hive> Describe RetailDB_EC2;
OK
event_time          string          from deserializer
product_id          string          from deserializer
category_id         string          from deserializer
category_code       string          from deserializer
brand               string          from deserializer
price               string          from deserializer
user_id             string          from deserializer
user_session        string          from deserializer
event_type          string

# Partition Information
# col_name          data_type      comment

event_type          string
Time taken: 0.05 seconds, Fetched: 14 row(s)
hive>

```

- Load the Data into the tables:

#### Query:

Insert into table RetailDB\_EU1 PARTITION (event\_type) select cast (replace (event\_time, 'UTC', '') as timestamp), product\_id, category\_id, category\_code, brand, cast(price as float), cast(user\_id as bigint), user\_session, event\_type from RetailDB;

Insert into table RetailDB\_EC2 PARTITION (event\_type) select cast (replace (event\_time, 'UTC', '') as timestamp), product\_id, category\_id, category\_code, brand, cast(price as float), cast(user\_id as bigint), user\_session, event\_type from RetailDB;

```

hive> Insert into table RetailDB_EU1 PARTITION (event_type) select cast (replace (event_time,
'UTC', '') as timestamp), product_id, category_id, category_code, brand, cast(price as float
), cast(user_id as bigint), user_session, event_type from RetailDB;
Query ID = hadoop_20220202073028_4d072466-47b9-469e-a749-129e2ec4e3d7
Total jobs = 1
Launching Job 1 out of 1

```

```
Loading data to table hive_data.retaildb_eul partition (event_type=null)
```

```
Time taken to load dynamic partitions: 0.683 seconds
Time taken for adding to write entity : 0.003 seconds
```

```
OK
```

```
Time taken: 243.338 seconds
```

```
hive> Insert into table RetailDB_EC2 PARTITION (event_type) select cast (replace (event_time,
'UTC', '') as timestamp), product_id, category_id, category_code, brand, cast(price as float
), cast(user_id as bigint), user_session, event_type from RetailDB;Insert into table RetailDB
_EC2 PARTITION (event_type) select cast (replace (event_time, 'UTC', '') as timestamp), produ
ct_id, category_id, category_code, brand, cast(price as float), cast(user_id as bigint), user
_session, event_type from RetailDB;
```

```
Query ID = hadoop_20220202073716_d0e603eb-acbc-4a00-b907-83b86c0f64e8
```

```
Total jobs = 1
```

```
Launching Job 1 out of 1
```

```
Status: Running (Executing on YARN cluster with App id application_1643783146095_0005)
```

```
Loading data to table hive_data.retaildb_ec2 partition (event_type=null)
```

```
Time taken to load dynamic partitions: 0.675 seconds
Time taken for adding to write entity : 0.001 seconds
```

```
OK
```

```
Time taken: 199.337 seconds
```

Further, we did run some basic query in both and found the second option to return the output in lesser time. Hence, we will continue with only second option.

### Query:

```
SELECT Avg(price) AS Oct_Revenue FROM RetailDB_EU1 WHERE MONTH(event_time) = '10' AND event_type =
'purchase';
```

```
hive> SELECT Avg(price) AS Oct_Revenue FROM RetailDB_EU1 WHERE MONTH(event_time) = '10' AND e
vent_type = 'purchase';
```

```
Query ID = hadoop_20220202074700_89c1b0bc-7b9a-4fb2-b486-63514055a70f
```

```
Total jobs = 1
```

```
Launching Job 1 out of 1
```

```
Status: Running (Executing on YARN cluster with App id application_1643783146095_0005)
```

```
Map 1: 0/3      Reducer 2: 0/1
Map 1: 0/3      Reducer 2: 0/1
Map 1: 0/3      Reducer 2: 0/1
Map 1: 0(+1)/3  Reducer 2: 0/1
Map 1: 0(+2)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 1(+2)/3  Reducer 2: 0/1
Map 1: 1(+2)/3  Reducer 2: 0(+1)/1
Map 1: 2(+1)/3  Reducer 2: 0(+1)/1
Map 1: 3/3      Reducer 2: 0(+1)/1
Map 1: 3/3      Reducer 2: 1/1
```

```
OK
```

```
4.93249206103583
```

```
Time taken: 29.227 seconds, Fetched: 1 row(s)
```

### Query:

```
SELECT Avg(price) AS Oct_Revenue FROM RetailDB_EC2 WHERE MONTH(event_time) = '10' AND event_type =
'purchase';
```

```

hive> SELECT Avg(price) AS Oct_Revenue FROM RetailDB_EC2 WHERE MONTH(event_time) = '10' AND e
vent_type = 'purchase';
Query ID = hadoop_20220202074808_7dc9bele-8530-447a-8a09-4b5dcd41af05
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1643783146095_0005)

Map 1: 0/4      Reducer 2: 0/1
Map 1: 0/4      Reducer 2: 0/1
Map 1: 0/4      Reducer 2: 0/1
Map 1: 0(+1)/4  Reducer 2: 0/1
Map 1: 0(+2)/4  Reducer 2: 0/1
Map 1: 0(+3)/4  Reducer 2: 0/1
Map 1: 0(+3)/4  Reducer 2: 0/1
Map 1: 0(+3)/4  Reducer 2: 0/1
Map 1: 0(+3)/4  Reducer 2: 0/1
Map 1: 0(+3)/4  Reducer 2: 0/1
Map 1: 0(+3)/4  Reducer 2: 0/1
Map 1: 0(+3)/4  Reducer 2: 0/1
Map 1: 1(+2)/4  Reducer 2: 0/1
Map 1: 3(+0)/4  Reducer 2: 0/1
Map 1: 3(+1)/4  Reducer 2: 0(+1)/1
Map 1: 4/4      Reducer 2: 0(+1)/1
Map 1: 4/4      Reducer 2: 1/1
OK
4.932492061034771
Time taken: 31.751 seconds, Fetched: 1 row(s)
hive>

```

So, we see, that time taken by RetailDB\_EU1 is slightly lesser than RetailDB\_EC2.

Hence, we will go ahead with RetailDB\_EU1 table for our further analysis:

## Questions and Queries to Answers:

Q1. Find the total revenue generated due to purchases made in October.

### Query:

```
SELECT sum(price)
FROM RetailDB_EU1
WHERE month(event_time)=10 and event_type='purchase';
```

### Answer:

```
hive> SELECT SUM(price) AS Oct_Total_Revenue FROM RetailDB_EU1 WHERE MONTH(event_time) = '10'
AND event_type = 'purchase';
Query ID = hadoop_20220202075907_e659172a-56da-4350-830f-5814530b21f0
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1643783146095_0006)

Map 1: 0/3      Reducer 2: 0/1
Map 1: 0/3      Reducer 2: 0/1
Map 1: 0/3      Reducer 2: 0/1
Map 1: 0(+2)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+2)/3  Reducer 2: 0/1
Map 1: 1(+2)/3  Reducer 2: 0(+1)/1
Map 1: 1(+2)/3  Reducer 2: 0(+1)/1
Map 1: 2(+1)/3  Reducer 2: 0(+1)/1
Map 1: 3/3      Reducer 2: 0(+1)/1
Map 1: 3/3      Reducer 2: 1/1
OK
1211538.4299998647
Time taken: 41.09 seconds, Fetched: 1 row(s)
```

Q2. Write a query to yield the total sum of purchases per month in a single output.

**Query:**

```
SELECT MONTH(event_time) AS Month, COUNT(event_type) AS Purchases
FROM RetailDB_EU1
WHERE event_type = 'purchase'
GROUP BY MONTH(event_time);
```

**Answer:**

```
hive> SELECT MONTH(event_time) AS Month, COUNT(event_type) AS Purchases FROM RetailDB_EU1 WHERE event_type = 'purchase' GROUP BY MONTH(event_time);
Query ID = hadoop_20220202080105_6c282587-9f05-43ae-841b-6afe0d74b787
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1643783146095_0006)

Map 1: 0/3      Reducer 2: 0/1
Map 1: 0/3      Reducer 2: 0/1
Map 1: 0/3      Reducer 2: 0/1
Map 1: 0(+1)/3  Reducer 2: 0/1
Map 1: 0(+2)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 1(+2)/3  Reducer 2: 0(+1)/1
Map 1: 2(+1)/3  Reducer 2: 0(+1)/1
Map 1: 3/3      Reducer 2: 0(+1)/1
Map 1: 3/3      Reducer 2: 1/1
OK
10      245624
11      322417
Time taken: 27.566 seconds, Fetched: 2 row(s)
```

Q3. Write a query to find the change in revenue generated due to purchases from October to November.

Query:

```
SELECT sum (case when month(event_time)=10 then price else -1*price end)
as change_in_revenue

FROM RetailDB_EU1

WHERE month(event_time) in (10,11) and event_type='purchase';
```

Answer:

```
hive> select sum (case when month(event_time)=10 then price else -1*price end) as change_in_r
evenue from RetailDB_EU1 where month(event_time) in (10,11) and event_type='purchase';
Query ID = hadoop_20220202080258_4da31942-4a05-4a93-aced-4b395a377729
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1643783146095_0006)

Map 1: 0/3      Reducer 2: 0/1
Map 1: 0/3      Reducer 2: 0/1
Map 1: 0/3      Reducer 2: 0/1
Map 1: 0(+2)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 1(+2)/3  Reducer 2: 0(+1)/1
Map 1: 1(+2)/3  Reducer 2: 0(+1)/1
Map 1: 2(+1)/3  Reducer 2: 0(+1)/1
Map 1: 3/3      Reducer 2: 0(+1)/1
Map 1: 3/3      Reducer 2: 1/1
OK
-319478.47000001266
Time taken: 29.882 seconds, Fetched: 1 row(s)
```



Q4. Find distinct categories of products. Categories with null category code can be ignored.

**Query:**

```
SELECT distinct split(category_code,'\\\.')[0] as cat
FROM RetailDB_EU1
WHERE split(category_code,'\\\.')[0] <> '';
```

**Answer:**

```
hive> select distinct split(category_code,'\\\.')[0] as cat from RetailDB_EC2 where split(category_code,'\\\.')[0] <> '';
Query ID = hadoop_20220202080639_7e03a331-fb29-40fd-accb-06a1e8b610ea
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1643783146095_0006)

Map 1: 0/6      Reducer 2: 0/9
```

```
Map 1: 6/6      Reducer 2: 7(+2)/9
Map 1: 6/6      Reducer 2: 9/9
OK
sport
stationery
accessories
appliances
category_code
furniture
apparel
Time taken: 119.917 seconds, Fetched: 7 row(s)
```

Q5. Find the total number of products available under each category.

**Query:**

```
SELECT SPLIT(category_code,'\\\.')[0] AS cat, COUNT(product_id) AS
No_of_products
FROM RetailDB_EU1
WHERE SPLIT(category_code,'\\\.')[0] <> ''
GROUP BY SPLIT(category_code,'\\\.')[0]
ORDER BY No_of_products DESC;
```

**Answer:**

```
hive> SELECT SPLIT(category_code,'\\\.')[0] AS cat, COUNT(product_id) AS No_of_products FROM R
etailDB_EU1 WHERE SPLIT(category_code,'\\\.')[0] <> '' GROUP BY SPLIT(category_code,'\\\.')[0]
ORDER BY No_of_products DESC;
Query ID = hadoop_20220202080916_clcca5c0-65af-49a6-8baf-8e8fb45e2bbf
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1643783146095_0006)
Map 1: 0/6      Reducer 2: 0/5  Reducer 3: 0/1
```

```
Map 1: 6/6      Reducer 2: 5/5  Reducer 3: 1/1
OK
appliances      61736
stationery      26722
furniture       23604
apparel 18232
accessories     12929
sport          2
category_code   2
Time taken: 81.426 seconds, Fetched: 7 row(s)
```

Q6. Which brand had the maximum sales in October and November combined?

**QUERY:**

```
SELECT brand, SUM(price) AS Sales
FROM RetailDB_EU1
WHERE brand <>' ' AND event_type='purchase'
GROUP BY brand
ORDER BY Sales DESC LIMIT 1 ;
```

**Answer:**

```
hive> SELECT brand, SUM(price) AS Sales FROM RetailDB_EU1 WHERE brand <>' ' AND event_type='purchase' GROUP BY brand ORDER BY Sales DESC LIMIT 1 ;
Query ID = hadoop_20220202081148_9f6d9d3b-0ab9-4990-a24e-a8681d2da5d5
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1643783146095_0006)

Map 1: 0/3      Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0/3      Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0/3      Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0(+1)/3 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0(+3)/3 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0(+3)/3 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0(+3)/3 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0(+3)/3 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0(+3)/3 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 1(+2)/3 Reducer 2: 0(+1)/1 Reducer 3: 0/1
Map 1: 2(+1)/3 Reducer 2: 0(+1)/1 Reducer 3: 0/1
Map 1: 3/3      Reducer 2: 0(+1)/1 Reducer 3: 0/1
Map 1: 3/3      Reducer 2: 1/1 Reducer 3: 0/1
Map 1: 3/3      Reducer 2: 1/1 Reducer 3: 0(+1)/1
Map 1: 3/3      Reducer 2: 1/1 Reducer 3: 1/1
OK
runail 148297.9400000006
Time taken: 27.715 seconds, Fetched: 1 row(s)
```

Q7. Which brands increased their sales from October to November?

**QUERY:**

WITH Brand\_Sales AS

```
(SELECT brand, round(SUM(CASE WHEN MONTH(event_time)=10 THEN price  
ELSE 0 END),2) AS October_sales, round(SUM(CASE WHEN  
MONTH(event_time)=11 THEN PRICE ELSE 0 END),2) AS November_sales  
FROM RetailDB_EU1  
WHERE event_type = 'purchase' AND MONTH(event_time) in ('10','11')  
AND brand != ''  
GROUP BY brand)
```

```
SELECT brand, October_sales, November_sales, round((November_sales -  
October_sales),2)AS Sales
```

```
FROM Brand_Sales
```

```
WHERE November_sales -October_sales > 0
```

```
ORDER BY Sales DESC;
```


**Answer:**

```

hadoop@ip-172-31-44-242:~
hive> WITH Brand_Sales AS(SELECT brand, round(SUM(CASE WHEN MONTH(event_time)=10 THEN price ELSE 0 END),2) AS October_sales,round(SUM(CASE WHEN MONTH(event_time)=11 THEN PRICE ELSE 0 END),2) AS November_sales FROM RetailDB_EU1 WHERE event_type = 'purchase' AND MONTH(event_time) in ('10','11') AND brand != '' GROUP BY brand) SELECT brand, October_sales, November_sales, round((November_sales -October_sales),2)AS Sales FROM Brand_Sales WHERE November_sales -October_sales > 0 ORDER BY Sales DESC;
Query ID = hadoop_20220202081307_f278a30a-f3a3-40f6-b6d8-422597c669ef
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1643783146095_0006)

Map 1: 0/3      Reducer 2: 0/1  Reducer 3: 0/1
Map 1: 0/3      Reducer 2: 0/1  Reducer 3: 0/1
Map 1: 0/3      Reducer 2: 0/1  Reducer 3: 0/1
Map 1: 0(+1)/3  Reducer 2: 0/1  Reducer 3: 0/1
Map 1: 0(+2)/3  Reducer 2: 0/1  Reducer 3: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1  Reducer 3: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1  Reducer 3: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1  Reducer 3: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1  Reducer 3: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1  Reducer 3: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1  Reducer 3: 0/1
Map 1: 1(+2)/3  Reducer 2: 0/1  Reducer 3: 0/1
Map 1: 1(+2)/3  Reducer 2: 0(+1)/1  Reducer 3: 0/1
Map 1: 1(+2)/3  Reducer 2: 0(+1)/1  Reducer 3: 0/1
Map 1: 2(+1)/3  Reducer 2: 0(+1)/1  Reducer 3: 0/1
Map 1: 3/3      Reducer 2: 0(+1)/1  Reducer 3: 0/1
Map 1: 3/3      Reducer 2: 1/1  Reducer 3: 0(+1)/1
Map 1: 3/3      Reducer 2: 1/1  Reducer 3: 1/1
OK
grattol 35445.54      71472.71      36027.17
uno      35302.03      51039.75      15737.72
lianail 5892.84 16394.24      10501.4
ingarden 23161.39      33566.21      10404.82
strong 29196.63      38671.27      9474.64
jessnail 26287.84      33345.23      7057.39
cosmoprofi 8322.81 14536.99      6214.18
polarus 6013.72 11371.93      5358.21
runail 71539.28      76758.66      5219.38
freedecor 3421.78 7671.8 4250.02
staleks 8519.73 11875.61      3355.88
bpw.style 11572.15      14837.44      3265.29
lovely 8704.38 11939.06      3234.68


```

 hadoop@ip-172-31-44-242:~

```

lovely 8704.38 11939.06 3234.68
marathon 7280.75 10273.1 2992.35
haruyama 9390.69 12352.91 2962.22
yoko 8756.91 11707.88 2950.97
italwax 21940.24 24799.37 2859.13
benovy 409.62 3259.97 2850.35
kaypro 881.34 3268.7 2387.36
estel 21756.75 24142.67 2385.92
concept 11032.14 13380.4 2348.26
kapous 11927.16 14093.08 2165.92
f.o.x 6624.23 8577.28 1953.05
masura 31266.08 33058.47 1792.39
milv 3904.94 5642.01 1737.07
beautix 10493.95 12222.95 1729.0
artex 2730.64 4327.25 1596.61
domix 10472.05 12009.17 1537.12
shik 3341.2 4839.72 1498.52
smart 4457.26 5902.14 1444.88
roubloff 3491.36 4913.77 1422.41
levrana 2243.56 3664.1 1420.54
oniq 8425.41 9841.65 1416.24
irisk 45591.96 46946.04 1354.08
severina 4775.88 6120.48 1344.6
joico 705.52 2015.1 1309.58
zeitun 708.66 2009.63 1300.97
beauty-free 554.17 1782.86 1228.69
swarovski 1887.93 3043.16 1155.23
de.lux 1659.7 2775.51 1115.81
metzger 5373.45 6457.16 1083.71
markell 1768.75 2834.43 1065.68
sanoto 157.14 1209.68 1052.54
nagaraku 4369.74 5327.68 957.94
ecolab 262.85 1214.3 951.45
art-visage 2092.71 2997.8 905.09
levissime 2227.5 3085.31 857.81
missha 1293.83 2150.28 856.45
solomeya 1899.7 2685.8 786.1
rosi 3077.04 3841.56 764.52
refectocil 2716.18 3475.58 759.4
kaaral 4412.43 5086.07 673.64
kosmekka 1181.44 1813.37 631.93
kinetics 6334.25 6945.26 611.01
browxenna 14331.37 14916.73 585.36


```

 hadoop@ip-172-31-44-242:~

```

browxenna      14331.37      14916.73      585.36
airnails       5118.9  5691.52  572.62
uskusi  5142.27  5690.31  548.04
coifin  903.0  1428.49  525.49
s.care  412.68  913.07  500.39
limoni  1308.9  1796.6  487.7
matrix  3243.25  3726.74  483.49
gehwol  1089.07  1557.68  468.61
greymy  29.21  489.49  460.28
bioaqua  942.89  1398.12  455.23
farmavita      837.37  1291.97  454.6
sophin  1067.86  1515.52  447.66
yu-r    271.41  673.71  402.3
kiss    421.55  817.33  395.78
naomi   0.0  389.0  389.0
lador   2083.61  2471.53  387.92
ellips  245.85  606.04  360.19
jas     3318.96  3657.43  338.47
lowence 242.84  567.75  324.91
nitrile 847.28  1162.68  315.4
shary   871.96  1176.49  304.53
kims    330.04  632.04  302.0
happyfons      801.92  1091.59  289.67
kocostar      310.85  594.93  284.08
insight 1443.7  1721.96  278.26
candy    534.96  799.38  264.42
bluesky 10307.24      10565.53      258.29
beauugreen    511.51  768.35  256.84
protokeratin  201.25  456.79  255.54
trind  298.07  542.96  244.89
entity  479.71  719.26  239.55
skinlite      651.94  890.45  238.51
provoc  827.99  1063.82  235.83
fedua   52.38  263.81  211.43
ecocraft      41.16  241.95  200.79
keen    236.35  435.62  199.27
mane    66.79  260.26  193.47
freshbubble   318.7  502.34  183.64
matreshka     0.0  182.67  182.67
chi    358.94  538.61  179.67
cristalinas   427.63  584.95  157.32
farmona 1692.46  1843.43  150.97
latinoil      249.52  384.59  135.07

```

 hadoop@ip-172-31-44-242:~

```

latinoil      249.52  384.59  135.07
miskin 158.04  293.07  135.03
elizavecca   70.53   204.3   133.77
nefertiti    233.52  366.64  133.12
finish  98.38   230.38  132.0
igrobeauty   513.66  645.07  131.41
dizao   819.13  945.51  126.38
osmo    645.58  762.31  116.73
batiste  772.4   874.17  101.77
carmex   145.08  243.36  98.28
eos      54.34   152.61  98.27
depilflax   2707.07  2803.78  96.71
enjoy    41.35   136.57  95.22
kerasys  430.91  525.2   94.29
aura     83.95   177.51  93.56
plazan   101.37  194.01  92.64
koelf    422.73  507.29  84.56
nirvel   163.04  234.33  71.29
konad    739.83  810.67  70.84
egomania   77.47   146.04  68.57
cutrin    299.37  367.62  68.25
laboratorium 246.5   312.52  66.02
inm      288.02  351.21  63.19
dewal     0.0    61.29   61.29
marutaka-foot 49.22  109.33  60.11
kares     0.0    59.45   59.45
profhenna  679.23  736.85  57.62
koelcia   55.5    112.75  57.25
balbcare   155.33  212.38  57.05
elskin    251.09  307.65  56.56
foamie    35.04   80.49   45.45
ladykin   125.65  170.57  44.92
likato    296.06  340.97  44.91
mavala    409.04  446.32  37.28
vilenta   197.6   231.21  33.61
beautyblender 78.74  109.41  30.67
biore     60.65   90.31   29.66
orly      902.38  931.09  28.71
estelare   444.81  471.87  27.06
profepil   93.36   118.02  24.66
blixz     38.95   63.4    24.45
binacil    0.0    24.26   24.26
godefroy   401.22  425.12  23.9

```



```
godefroy      401.22  425.12  23.9
glysolid      69.73   91.59  21.86
veraclara     50.11   71.21  21.1
juno          0.0    21.08  21.08
kamill 63.01   81.49  18.48
treaclemoon   163.37  181.49  18.12
supertan      50.37   66.51  16.14
barbie 0.0     12.39  12.39
deoproce      316.84  329.17  12.33
rasyan 18.8    28.94  10.14
fly 17.14     27.17  10.03
tertio 236.16  245.8   9.64
jaguar 1102.11  1110.65  8.54
soleo 204.2    212.53  8.33
neoleor 43.41   51.7    8.29
moyou 5.71     10.28  4.57
bodyton 1376.34  1380.64  4.3
skinity 8.88    12.44  3.56
helloganic    0.0     3.1     3.1
grace 100.92   102.61  1.69
cosima 20.23   20.93  0.7
ovale 2.54     3.1     0.56
Time taken: 30.313 seconds, Fetched: 160 row(s)
hive>
```

Q8. Your company wants to reward the top 10 users of its website with a Golden Customer plan. Write a query to generate a list of top 10 users who spend the most.

**Query:**

```
SELECT user_id, SUM(price) AS Total_Expense
FROM RetailDB_EU1
WHERE event_type='purchase'
GROUP BY user_id
ORDER BY Total_Expense DESC LIMIT 10 ;
```

**Answer:**

```
hive> SELECT user_id, SUM(price) AS Total_Expense FROM RetailDB_EC2 WHERE event_type='purchase' GROUP BY user_id ORDER BY Total_Expense DESC LIMIT 10 ;
Query ID = hadoop_20220202081746_b7ca436f-4bf6-4f65-965a-c17c0936e9e8
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1643783146095_0006)

Map 1: 0/4      Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0/4      Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0/4      Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0(+1)/4 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0(+3)/4 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0(+3)/4 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0(+3)/4 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0(+3)/4 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0(+3)/4 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0(+3)/4 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 0(+3)/4 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 1(+2)/4 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 1(+3)/4 Reducer 2: 0/1 Reducer 3: 0/1
Map 1: 2(+2)/4 Reducer 2: 0(+1)/1 Reducer 3: 0/1
Map 1: 3(+1)/4 Reducer 2: 0(+1)/1 Reducer 3: 0/1
Map 1: 4/4      Reducer 2: 0(+1)/1 Reducer 3: 0/1
Map 1: 4/4      Reducer 2: 1/1 Reducer 3: 0(+1)/1
Map 1: 4/4      Reducer 2: 1/1 Reducer 3: 1/1
OK
557790271      5431.7399999999985
150318419      3291.9399999999982
562167663      2705.7
531900924      2658.9
557850743      2590.9599999999996
522130011      2370.7799999999984
561592095      2219.4
431950134      2195.1799999999994
566576008      2112.7200000000025
521347209      2081.82
Time taken: 32.848 seconds, Fetched: 10 row(s)
```

## Dropping the Database and Terminating the Cluster

### Drop Database:

- Run the below queries to check out database:

#### Query:

Show databases;

```
hive> show databases;
OK
default
hive_data
Time taken: 0.011 seconds, Fetched: 2 row(s)
```

Drop the “hive\_data” database (Use cascade to drop the tables inside before dropping the database):

#### Query:

Drop database hive\_data cascade;

```
hive> drop database hive_data cascade;
OK
Time taken: 0.499 seconds
```

- Reconfirm if database is dropped using (Show databases;) :

```
hive> show databases;
OK
default
Time taken: 0.011 seconds, Fetched: 1 row(s)
hive> █
```

- Quit Hive:

```
hive> Quit
> ;
[hadoop@ip-172-31-44-242 ~]$ █
```

- Type “Exit” to quit from putty.

### Terminate EMR:

- Select Terminate > Select Terminate in the pop-up window >Status changes to Terminating.

### Terminate cluster

Are you sure you want to terminate this cluster?

Any pending work or data residing on the cluster will be lost, such as data stored in HDFS. This action is irreversible.

[Cancel](#)[Terminate](#)

Create cluster

View details

Clone

Terminate

Filter:

All clusters

Filter clusters ...

18 clusters (all loaded)

	Name	ID	Status	Creation time (UTC+5:30)	Elapsed time
<div><div></div><div></div></div>	<a href="#">Hive Case Study</a>	j-2KSY1QFM8Q2KW	Terminated User request	2022-02-02 11:47 (UTC+5:30)	2 hours, 19 minutes

**Cluster is Terminated!**